

**МИНИСТЕРСТВО НАУКИ И ВЫСШЕГО ОБРАЗОВАНИЯ
РОССИЙСКОЙ ФЕДЕРАЦИИ
ФЕДЕРАЛЬНОЕ ГОСУДАРСТВЕННОЕ БЮДЖЕТНОЕ ОБРАЗОВАТЕЛЬНОЕ УЧРЕЖДЕНИЕ
ВЫСШЕГО ОБРАЗОВАНИЯ
«БЕЛГОРОДСКИЙ ГОСУДАРСТВЕННЫЙ
ТЕХНОЛОГИЧЕСКИЙ УНИВЕРСИТЕТ им. В. Г. ШУХОВА»
(БГТУ им. В.Г. Шухова)**



ИНСТИТУТ ИНФОРМАЦИОННЫХ ТЕХНОЛОГИЙ И УПРАВЛЯЮЩИХ СИСТЕМ

Лабораторная работа №8

по дисциплине: Базы данных

тема: «Оптимизация SQL-запросов в СУБД с использованием планировщика»

Выполнил: ст. группы ПВ-223
Пахомов Владислав Андреевич

Проверили:
ст. пр. Панченко Максим Владимирович

Белгород 2024 г.

Лабораторная работа №8

Оптимизация SQL-запросов в СУБД с использованием планировщика

Вариант 8

Цель работы: получение навыков повышения производительности работы СУБД с помощью оптимизации sql-запросов с использованием планировщика.

Выполним генерацию данных при помощи faker, дополнив прошлую программу:

fake.py

```
import sys

from PySide6.QtWidgets import QApplication
from sqlalchemy.orm import Session

from core.db import engine
from models import Home, Contract, Resident, Payment, Task, Worker
from faker import Faker
from datetime import date, timedelta
from widgets.main import MainDialog
import random

MAX_COUNT = 100

homes = []
contracts = []
residents = []
payments = []
tasks = []
workers = []

if __name__ == "__main__":
    fake = Faker()
    with Session(engine) as session:
        print("Generating homes...")
        for i in range(0, MAX_COUNT):
            new_home = Home(
                address=fake.address(),
                commissioning=fake.date_between(
                    start_date=(date.today() + timedelta(days=500)),
                    end_date=(date.today() + timedelta(days=1000))),
                floors=fake.random_int(1, 50),
                index=fake.random_int(100000, 999999)
            )

            session.add(new_home)
            session.commit()
            homes.append(new_home)

        print("Generating contracts...")
        for i in range(0, MAX_COUNT):
            new_contract = Contract(
                transaction_date=fake.date_between(
                    start_date=(date.today() + timedelta(days=500)),
                    end_date=(date.today() + timedelta(days=1000))),
                until_date=fake.date_between(
                    start_date=(date.today() + timedelta(days=1001)),
                    end_date=(date.today() + timedelta(days=1500))),
                home_address=random.choice(homes).address
            )

            session.add(new_contract)
            session.commit()
```

```

        contracts.append(new_contract)

    prefix_phone = [
        "+7921",
        "+7931",
        "+7911",
        "+7953",
        "+7995",
        "+7912",
        "+7991",
        "+7999",
        "+7956",
        "+7900",
    ]

    print("Generating residents...")
    for i in range(0, MAX_COUNT):
        new_resident = Resident(
            passport_data=str(fake.random_int(1000_000000, 9999_999999)),
            surname=fake.last_name(),
            name=fake.first_name(),
            patronymics=fake.first_name() if random.random() > 0.5 else None,
            email=fake.email(),
            phone=f"{random.choice(prefix_phone)}{random.randint(100,
999)}{random.randint(1000, 9999)}",
            contracts=random.sample(contracts, fake.random_int(0, 2))
        )

        session.add(new_resident)
        session.commit()
        residents.append(new_resident)

    energy_sources = [
        "Ремонт",
        "Интернет",
        "Вода",
        "Электричество",
        "Отопление"
    ]

    print("Generating payments...")
    for i in range(0, MAX_COUNT):
        new_payment = Payment(
            id=str(fake.random_int(1000000000000, 10000000000000)),
            until_date=fake.date_between(
                start_date=(date.today() - timedelta(days=300)),
                end_date=(date.today() - timedelta(days=200))),
            paid_date=fake.date_between(
                start_date=(date.today() - timedelta(days=400)),
                end_date=(date.today() - timedelta(days=100))) if random.random() > 0.5 else
None,
            contract_id=random.choice(contracts).id,
            energy_source=random.choice(energy_sources),
            payment=fake.random_int(100, 10001)
        )

        session.add(new_payment)
        session.commit()
        payments.append(new_payment)

    print("Generating tasks...")
    for i in range(0, MAX_COUNT):
        new_task = Task(
            payment=fake.random_int(100_000, 10_000_000),
            until_date=fake.date_between(
                start_date=(date.today() - timedelta(days=300)),
                end_date=(date.today() - timedelta(days=200))),

```

```

        completed_date=fake.date_between(
            start_date=(date.today() - timedelta(days=400)),
            end_date=(date.today() - timedelta(days=100))) if random.random() > 0.5 else
None,

        home_address=random.choice(homes).address
    )

    session.add(new_task)
    session.commit()
    tasks.append(new_task)

print("Generating tasks...")
for i in range(0, MAX_COUNT):
    new_worker = Worker(
        inn=str(fake.random_int(100_000_000_000, 1_000_000_000_000)),
        email=fake.email() if random.random() < 0.5 else None,
        phone=f"{random.choice(prefix_phone)}{random.randint(100,
999)}{random.randint(1000, 9999)}" if random.random() < 0.5 else None,
        tasks=random.sample(tasks, fake.random_int(0, 10))
    )

    session.add(new_worker)
    session.commit()
    workers.append(new_worker)

print("Done!")

```

run_fake.ps1

```

echo "> Read .env file"

# Default location of venv file
set-content env:\VENV_DIR .venv

get-content .env | foreach {
    $name, $value = $_.split('=')
    set-content env:\$name $value
}

if (!(Test-Path "$env:VENV_DIR/Scripts/python.exe")){
    echo "> Venv is not found, let's install it. It'd take some time, grab some coffee!"
    if (Test-Path "$env:VENV_DIR") {
        rm -r "$env:VENV_DIR"
    }
    mkdir "$env:VENV_DIR"
    python -m venv "$env:VENV_DIR"
}

echo "> Installing requirments"
powershell "$env:VENV_DIR/Scripts/pip.exe install -r ./requirments.txt"

powershell "$env:VENV_DIR/Scripts/alembic-autogen-check"
if (!$?) {
    echo "> Database is not up-to-date"
    echo "> Autogen new revision"
    powershell "$env:VENV_DIR/Scripts/alembic revision --autogenerate"
    echo "> Upgrade to latest alembic version"
    powershell "$env:VENV_DIR/Scripts/alembic upgrade head"
} else {
    echo "> Database is up-to-date"
}

echo "> Start application"
powershell "$env:VENV_DIR/Scripts/python.exe fake.py"

```

До оптимизаций:

EXPLAIN ANALYZE of non payers

```
Sort (cost=171.20..172.58 rows=550 width=68) (actual time=1.112..1.125 rows=481 loops=1)
  Sort Key: (sum(payment.payment))
  Sort Method: quicksort Memory: 64kB
  -> HashAggregate (cost=139.29..146.17 rows=550 width=68) (actual time=0.892..1.028 rows=481 loops=1)
    Group Key: resident.passport_data, payment.energy_source
    Batches: 1 Memory Usage: 105kB
    -> Hash Join (cost=100.54..135.17 rows=550 width=52) (actual time=0.536..0.749 rows=522 loops=1)
      Hash Cond: ((resident.passport_data)::text =
(residents_contracts.resident_passport_data)::text)
      -> Seq Scan on resident (cost=0.00..25.00 rows=1100 width=31) (actual time=0.003..0.109 rows=1100 loops=1)
      -> Hash (cost=93.67..93.67 rows=550 width=32) (actual time=0.528..0.528 rows=522 loops=1)
        Buckets: 1024 Batches: 1 Memory Usage: 42kB
        -> Hash Join (cost=67.14..93.67 rows=550 width=32) (actual time=0.335..0.471 rows=522 loops=1)
          Hash Cond: (residents_contracts.contract_id = contract.id)
          -> Seq Scan on residents_contracts (cost=0.00..16.93 rows=1093 width=15) (actual time=0.001..0.036 rows=1093 loops=1)
          -> Hash (cost=60.22..60.22 rows=554 width=29) (actual time=0.332..0.332 rows=554 loops=1)
            Buckets: 1024 Batches: 1 Memory Usage: 43kB
            -> Hash Join (cost=36.75..60.22 rows=554 width=29) (actual time=0.130..0.265 rows=554 loops=1)
              Hash Cond: (payment.contract_id = contract.id)
              -> Seq Scan on payment (cost=0.00..22.00 rows=554 width=25) (actual time=0.002..0.075 rows=554 loops=1)
              Filter: (paid_date IS NULL)
              Rows Removed by Filter: 546
              -> Hash (cost=23.00..23.00 rows=1100 width=4) (actual time=0.126..0.126 rows=1100 loops=1)
                Buckets: 2048 Batches: 1 Memory Usage: 55kB
                -> Seq Scan on contract (cost=0.00..23.00 rows=1100 width=4) (actual time=0.001..0.055 rows=1100 loops=1)
```

Planning Time: 0.259 ms

Execution Time: 1.164 ms

EXPLAIN ANALYZE of workers rating

```
Sort (cost=488.26..490.74 rows=992 width=53) (actual time=3.593..3.620 rows=992 loops=1)
  Sort Key: (COALESCE(anon_1.completed, '0'::bigint)) DESC
  Sort Method: quicksort Memory: 81kB
  -> Hash Left Join (cost=403.45..438.89 rows=992 width=53) (actual time=3.023..3.492 rows=992 loops=1)
    Hash Cond: ((worker.inn)::text = (anon_1.worker_inn)::text)
    -> Hash Join (cost=212.71..235.61 rows=992 width=21) (actual time=1.739..1.916 rows=992 loops=1)
      Hash Cond: ((worker.inn)::text = (anon_2.worker_inn)::text)
      -> Seq Scan on worker (cost=0.00..20.00 rows=1100 width=13) (actual time=0.002..0.041 rows=1100 loops=1)
      -> Hash (cost=200.31..200.31 rows=992 width=21) (actual time=1.736..1.737 rows=992 loops=1)
        Buckets: 1024 Batches: 1 Memory Usage: 60kB
        -> Subquery Scan on anon_2 (cost=180.47..200.31 rows=992 width=21) (actual time=1.523..1.654 rows=992 loops=1)
          -> HashAggregate (cost=180.47..190.39 rows=992 width=21) (actual time=1.523..1.601 rows=992 loops=1)
            Group Key: workers_tasks.worker_inn
            Batches: 1 Memory Usage: 129kB
            -> Hash Join (cost=43.25..152.05 rows=5683 width=13) (actual time=0.152..0.902 rows=5683 loops=1)
              Hash Cond: (workers_tasks.task_id = task.id)
```

```

        -> Seq Scan on workers_tasks (cost=0.00..93.83 rows=5683
width=17) (actual time=0.001..0.192 rows=5683 loops=1)
        -> Hash (cost=29.50..29.50 rows=1100 width=4) (actual
time=0.149..0.149 rows=1100 loops=1)
                Buckets: 2048 Batches: 1 Memory Usage: 55kB
        -> Seq Scan on task (cost=0.00..29.50 rows=1100
width=4) (actual time=0.002..0.090 rows=1100 loops=1)
                Filter: ((until_date > '2004-01-01'::date) AND
(until_date < '2040-01-01'::date))
        -> Hash (cost=178.35..178.35 rows=992 width=21) (actual time=1.279..1.280 rows=885
loops=1)
                Buckets: 1024 Batches: 1 Memory Usage: 54kB
        -> Subquery Scan on anon_1 (cost=158.51..178.35 rows=992 width=21) (actual
time=1.073..1.198 rows=885 loops=1)
                -> HashAggregate (cost=158.51..168.43 rows=992 width=21) (actual
time=1.073..1.144 rows=885 loops=1)
                        Group Key: workers_tasks_1.worker_inn
                        Batches: 1 Memory Usage: 129kB
                        -> Hash Join (cost=36.09..144.89 rows=2723 width=13) (actual
time=0.100..0.719 rows=2704 loops=1)
                                Hash Cond: (workers_tasks_1.task_id = task_1.id)
                                -> Seq Scan on workers_tasks workers_tasks_1 (cost=0.00..93.83
rows=5683 width=17) (actual time=0.001..0.183 rows=5683 loops=1)
                                -> Hash (cost=29.50..29.50 rows=527 width=4) (actual
time=0.098..0.098 rows=527 loops=1)
                                        Buckets: 1024 Batches: 1 Memory Usage: 27kB
                                -> Seq Scan on task task_1 (cost=0.00..29.50 rows=527
width=4) (actual time=0.001..0.069 rows=527 loops=1)
                                        Filter: ((completed_date IS NOT NULL) AND (until_date
> '2004-01-01'::date) AND (until_date < '2040-01-01'::date))
                                        Rows Removed by Filter: 573
Planning Time: 0.181 ms
Execution Time: 3.667 ms
EXPLAIN ANALYZE of house profit
Sort (cost=228.87..231.62 rows=1100 width=52) (actual time=1.546..1.576 rows=1100 loops=1)
  Sort Key: ((COALESCE(anon_1.plus, '0'::bigint) - COALESCE(anon_2.minus, '0'::bigint)))
  Sort Method: quicksort Memory: 134kB
  -> Hash Left Join (cost=141.76..173.31 rows=1100 width=52) (actual time=1.067..1.417
rows=1100 loops=1)
        Hash Cond: ((home.address)::text = (anon_2.home)::text)
        -> Hash Left Join (cost=89.93..115.83 rows=1100 width=52) (actual time=0.675..0.862
rows=1100 loops=1)
                Hash Cond: ((home.address)::text = (anon_1.home)::text)
                -> Seq Scan on home (cost=0.00..23.00 rows=1100 width=44) (actual
time=0.002..0.041 rows=1100 loops=1)
                -> Hash (cost=81.17..81.17 rows=701 width=52) (actual time=0.672..0.673 rows=519
loops=1)
                        Buckets: 1024 Batches: 1 Memory Usage: 51kB
                        -> Subquery Scan on anon_1 (cost=67.15..81.17 rows=701 width=52) (actual
time=0.545..0.619 rows=519 loops=1)
                                -> HashAggregate (cost=67.15..74.16 rows=701 width=52) (actual
time=0.545..0.589 rows=519 loops=1)
                                        Group Key: contract.home_address
                                        Batches: 1 Memory Usage: 105kB
                                        -> Hash Left Join (cost=36.75..61.65 rows=1100 width=48)
(actual time=0.184..0.385 rows=1100 loops=1)
                                                Hash Cond: (payment.contract_id = contract.id)
                                                -> Seq Scan on payment (cost=0.00..22.00 rows=1100
width=8) (actual time=0.002..0.084 rows=1100 loops=1)
                                                -> Hash (cost=23.00..23.00 rows=1100 width=48) (actual
time=0.180..0.180 rows=1100 loops=1)
                                                        Buckets: 2048 Batches: 1 Memory Usage: 104kB
                                                        -> Seq Scan on contract (cost=0.00..23.00 rows=1100
width=48) (actual time=0.002..0.072 rows=1100 loops=1)
                                                        -> Hash (cost=43.24..43.24 rows=687 width=52) (actual time=0.389..0.390 rows=687
loops=1)
                                                                Buckets: 1024 Batches: 1 Memory Usage: 65kB

```

```
-> Subquery Scan on anon_2 (cost=29.50..43.24 rows=687 width=52) (actual
time=0.230..0.322 rows=687 loops=1)
  -> HashAggregate (cost=29.50..36.37 rows=687 width=52) (actual
time=0.229..0.283 rows=687 loops=1)
    Group Key: task.home_address
    Batches: 1 Memory Usage: 169kB
    -> Seq Scan on task (cost=0.00..24.00 rows=1100 width=48) (actual
time=0.001..0.041 rows=1100 loops=1)
Planning Time: 0.129 ms
Execution Time: 1.626 ms
```

1. Рассмотрим запрос на получение прибыли домов. Можно попробовать оптимизировать запрос путём введения индексов на home_address:

```
EXPLAIN ANALYZE of house profit
Sort (cost=228.87..231.62 rows=1100 width=52) (actual time=1.667..1.698 rows=1100 loops=1)
  Sort Key: ((COALESCE(anon_1.plus, '0'::bigint) - COALESCE(anon_2.minus, '0'::bigint)))
  Sort Method: quicksort Memory: 134kB
  -> Hash Left Join (cost=141.76..173.31 rows=1100 width=52) (actual time=1.079..1.519
rows=1100 loops=1)
    Hash Cond: ((home.address)::text = (anon_2.home)::text)
    -> Hash Left Join (cost=89.93..115.83 rows=1100 width=52) (actual
time=0.674..0.876 rows=1100 loops=1)
      Hash Cond: ((home.address)::text = (anon_1.home)::text)
      -> Seq Scan on home (cost=0.00..23.00 rows=1100 width=44) (actual
time=0.002..0.043 rows=1100 loops=1)
      -> Hash (cost=81.17..81.17 rows=701 width=52) (actual time=0.670..0.671
rows=519 loops=1)
        Buckets: 1024 Batches: 1 Memory Usage: 51kB
        -> Subquery Scan on anon_1 (cost=67.15..81.17 rows=701 width=52)
(actual time=0.544..0.619 rows=519 loops=1)
          -> HashAggregate (cost=67.15..74.16 rows=701 width=52) (actual
time=0.544..0.589 rows=519 loops=1)
            Group Key: contract.home_address
            Batches: 1 Memory Usage: 105kB
            -> Hash Left Join (cost=36.75..61.65 rows=1100 width=48)
(actual time=0.180..0.381 rows=1100 loops=1)
              Hash Cond: (payment.contract_id = contract.id)
              -> Seq Scan on payment (cost=0.00..22.00 rows=1100
width=8) (actual time=0.002..0.082 rows=1100 loops=1)
              -> Hash (cost=23.00..23.00 rows=1100 width=48)
(actual time=0.176..0.177 rows=1100 loops=1)
                Buckets: 2048 Batches: 1 Memory Usage: 104kB
                -> Seq Scan on contract (cost=0.00..23.00
rows=1100 width=48) (actual time=0.001..0.071 rows=1100 loops=1)
            -> Hash (cost=43.24..43.24 rows=687 width=52) (actual time=0.402..0.403 rows=687
loops=1)
              Buckets: 1024 Batches: 1 Memory Usage: 65kB
              -> Subquery Scan on anon_2 (cost=29.50..43.24 rows=687 width=52) (actual
time=0.231..0.324 rows=687 loops=1)
                -> HashAggregate (cost=29.50..36.37 rows=687 width=52) (actual
time=0.230..0.284 rows=687 loops=1)
                  Group Key: task.home_address
                  Batches: 1 Memory Usage: 169kB
                  -> Seq Scan on task (cost=0.00..24.00 rows=1100 width=48)
(actual time=0.001..0.040 rows=1100 loops=1)
Planning Time: 0.126 ms
Execution Time: 1.761 ms
```

Добавленный индекс игнорируется планировщиком и только будет занимать время работы. Оптимизировать данный запрос не получится. Индекс не помог при оптимизации.

2. Попробуем оптимизировать запрос на получение рейтинга рабочих.

Основной проблемой этого запроса является то, что в нём используются два практически одинаковых подзапроса для вычисления `completed` и `total`.

```
completed = (  
    select(  
        WorkerTask.worker_inn,  
        func.count().label("completed")  
    )  
    .select_from(WorkerTask)  
    .join(Task)  
    .where(  
        and_(  
            Task.completed_date.isnot(None),  
            Task.until_date > begin_date,  
            Task.until_date < end_date  
        )  
    )  
    .group_by(WorkerTask.worker_inn)  
)  
.subquery()  
  
total = (  
    select(  
        WorkerTask.worker_inn,  
        func.count().label("total")  
    )  
    .select_from(WorkerTask)  
    .join(Task)  
    .where(  
        and_(  
            Task.until_date > begin_date,  
            Task.until_date < end_date  
        )  
    )  
    .group_by(WorkerTask.worker_inn)  
)  
.subquery()
```

Возможно, получится обойтись одним подзапросом:

```
completed_total = (  
    select(  
        WorkerTask.worker_inn,  
        func.count().label("total"),  
        func.count(Task.completed_date).label("completed")  
    )  
    .select_from(WorkerTask)  
    .join(Task)  
    .where(  
        Task.until_date.between(begin_date, end_date)  
    )  
    .group_by(WorkerTask.worker_inn)  
)  
.subquery()  
  
main_query = (select(  
    Worker.inn.label("worker_inn"),  
    func.coalesce(completed_total.c.completed, 0).label("completed"),  
    (1.0 * func.coalesce(completed_total.c.completed, 0) /  
    func.coalesce(completed_total.c.total, 1)).label("rating")  
    )  
    .select_from(Worker)  
    .join(completed_total, completed_total.c.worker_inn == Worker.inn, isouter=True)  
    .order_by(desc("completed")))
```



```

EXPLAIN ANALYZE of workers rating
Sort (cost=316.38..319.13 rows=1100 width=53) (actual time=2.426..2.454 rows=1100 loops=1)
  Sort Key: (COALESCE(anon_1.completed, '0'::bigint)) DESC
  Sort Method: quicksort Memory: 111kB
  -> Hash Left Join (cost=226.92..260.81 rows=1100 width=53) (actual time=1.944..2.323 rows=1100 loops=1)
    Hash Cond: ((worker.inn)::text = (anon_1.worker_inn)::text)
    -> Seq Scan on worker (cost=0.00..20.00 rows=1100 width=13) (actual time=0.002..0.039 rows=1100 loops=1)
    -> Hash (cost=214.52..214.52 rows=992 width=29) (actual time=1.938..1.939 rows=992 loops=1)
      Buckets: 1024 Batches: 1 Memory Usage: 68kB
      -> Subquery Scan on anon_1 (cost=194.68..214.52 rows=992 width=29) (actual time=1.706..1.852 rows=992 loops=1)
        -> HashAggregate (cost=194.68..204.60 rows=992 width=29) (actual time=1.706..1.795 rows=992 loops=1)
          Group Key: workers_tasks.worker_inn
          Batches: 1 Memory Usage: 193kB
          -> Hash Join (cost=43.25..152.05 rows=5683 width=17) (actual time=0.171..1.008 rows=5683 loops=1)
            Hash Cond: (workers_tasks.task_id = task.id)
            -> Seq Scan on workers_tasks (cost=0.00..93.83 rows=5683 width=17) (actual time=0.001..0.190 rows=5683 loops=1)
            -> Hash (cost=29.50..29.50 rows=1100 width=8) (actual time=0.169..0.169 rows=1100 loops=1)
              Buckets: 2048 Batches: 1 Memory Usage: 57kB
              -> Seq Scan on task (cost=0.00..29.50 rows=1100 width=8) (actual time=0.002..0.097 rows=1100 loops=1)
                Filter: ((until_date > '2004-01-01'::date) AND (until_date < '2040-01-01'::date))
Planning Time: 0.111 ms
Execution Time: 2.497 ms

```

Время подготовки и выполнения снизилось с

```

Planning Time: 0.181 ms
Execution Time: 3.667 ms

```

до

```

Planning Time: 0.111 ms
Execution Time: 2.497 ms

```

(практически в 1.5 раза)

А значит оптимизация прошла успешно. Можно также попробовать добавить индекс на поле `until_date`, чтобы сравнение с переданной датой выполнялось по индексу:

```

EXPLAIN ANALYZE of workers rating
Sort (cost=316.38..319.13 rows=1100 width=53) (actual time=2.376..2.404 rows=1100 loops=1)
  Sort Key: (COALESCE(anon_1.completed, '0'::bigint)) DESC
  Sort Method: quicksort Memory: 111kB
  -> Hash Left Join (cost=226.92..260.81 rows=1100 width=53) (actual time=1.894..2.272 rows=1100 loops=1)
    Hash Cond: ((worker.inn)::text = (anon_1.worker_inn)::text)
    -> Seq Scan on worker (cost=0.00..20.00 rows=1100 width=13) (actual time=0.002..0.038 rows=1100 loops=1)
    -> Hash (cost=214.52..214.52 rows=992 width=29) (actual time=1.889..1.890 rows=992 loops=1)
      Buckets: 1024 Batches: 1 Memory Usage: 68kB
      -> Subquery Scan on anon_1 (cost=194.68..214.52 rows=992 width=29) (actual time=1.658..1.803 rows=992 loops=1)
        -> HashAggregate (cost=194.68..204.60 rows=992 width=29) (actual time=1.657..1.746 rows=992 loops=1)
          Group Key: workers_tasks.worker_inn

```

```

      Batches: 1 Memory Usage: 193kB
      -> Hash Join (cost=43.25..152.05 rows=5683 width=17) (actual
time=0.181..1.010 rows=5683 loops=1)
        Hash Cond: (workers_tasks.task_id = task.id)
        -> Seq Scan on workers_tasks (cost=0.00..93.83 rows=5683
width=17) (actual time=0.001..0.189 rows=5683 loops=1)
        -> Hash (cost=29.50..29.50 rows=1100 width=8) (actual
time=0.178..0.178 rows=1100 loops=1)
          Buckets: 2048 Batches: 1 Memory Usage: 57kB
          -> Seq Scan on task (cost=0.00..29.50 rows=1100
width=8) (actual time=0.002..0.097 rows=1100 loops=1)
          Filter: ((until_date >= '2004-01-01'::date) AND
(until_date <= '2040-01-01'::date))
Planning Time: 0.123 ms
Execution Time: 2.448 ms

```

Однако индекс был проигнорирован планировщиком, время незначительно уменьшилось:

```

Planning Time: 0.123 ms
Execution Time: 2.448 ms

```

Оптимизация с применением индекса не принесла пользы.

3. Попробуем оптимизировать запрос на получение неплательщиков

Попробуем добавить индекс на Payment.paid_date, чтобы фильтр по запросу IS NULL выполнялся по индексированной таблице:

```

EXPLAIN ANALYZE of non payers
Sort (cost=171.20..172.58 rows=550 width=68) (actual time=1.099..1.112 rows=481 loops=1)
  Sort Key: (sum(payment.payment))
  Sort Method: quicksort Memory: 64kB
  -> HashAggregate (cost=139.29..146.17 rows=550 width=68) (actual time=0.880..1.016
rows=481 loops=1)
    Group Key: resident.passport_data, payment.energy_source
    Batches: 1 Memory Usage: 105kB
    -> Hash Join (cost=100.54..135.17 rows=550 width=52) (actual time=0.539..0.738
rows=522 loops=1)
      Hash Cond: ((resident.passport_data)::text =
(residents_contracts.resident_passport_data)::text)
      -> Seq Scan on resident (cost=0.00..25.00 rows=1100 width=31) (actual
time=0.003..0.097 rows=1100 loops=1)
      -> Hash (cost=93.67..93.67 rows=550 width=32) (actual time=0.530..0.530
rows=522 loops=1)
        Buckets: 1024 Batches: 1 Memory Usage: 42kB
        -> Hash Join (cost=67.14..93.67 rows=550 width=32) (actual
time=0.335..0.471 rows=522 loops=1)
          Hash Cond: (residents_contracts.contract_id = contract.id)
          -> Seq Scan on residents_contracts (cost=0.00..16.93 rows=1093
width=15) (actual time=0.001..0.036 rows=1093 loops=1)
          -> Hash (cost=60.22..60.22 rows=554 width=29) (actual
time=0.332..0.332 rows=554 loops=1)
            Buckets: 1024 Batches: 1 Memory Usage: 43kB
            -> Hash Join (cost=36.75..60.22 rows=554 width=29)
(actual time=0.130..0.274 rows=554 loops=1)
              Hash Cond: (payment.contract_id = contract.id)
              -> Seq Scan on payment (cost=0.00..22.00 rows=554
width=25) (actual time=0.002..0.074 rows=554 loops=1)
              Filter: (paid_date IS NULL)
              Rows Removed by Filter: 546
              -> Hash (cost=23.00..23.00 rows=1100 width=4)
(actual time=0.127..0.127 rows=1100 loops=1)
                Buckets: 2048 Batches: 1 Memory Usage: 55kB

```

```
-> Seq Scan on contract (cost=0.00..23.00  
rows=1100 width=4) (actual time=0.001..0.055 rows=1100 loops=1)  
Planning Time: 0.268 ms  
Execution Time: 1.152 ms
```

Индекс был проигнорирован, время изменилось незначительно

```
Planning Time: 0.268 ms  
Execution Time: 1.152 ms
```

Оптимизация с применением индекса не принесла пользы.

Ссылка на репозиторий: <https://github.com/IAmProgrammist/database/tree/main/lab8>

Вывод: в ходе лабораторной работы получили навыки повышения производительности работы СУБД с помощью оптимизации sql-запросов с использованием планировщика. Индексы, хоть и являются очень полезными, однако часто игнорируются планировщиками, основным инструментом при оптимизации запросов являются оптимальные запросы: использование как можно меньшего кол-ва вложенных запросов и джойнов, знание работы функций.